Breast Cancer Diagnosis Analysis

Predict breast cancer diagnosis (malignant vs benign) from digitized cell nuclei images using machin

Jotty SwarmMLComprehensive

February 05, 2026

Contents

Contents

0.1 Executive Summary

Predict breast cancer diagnosis (malignant vs benign) from digitized cell nuclei images using machine learning.

0.1.1 Key Results

Best Model: Logistic Regression

Performance Metrics:

Metric	Value
Accuracy	0.9825
Precision	0.9861
Recall	0.9861
F1	0.9861
Auc Roc	0.9954

Dataset: 30 features analyzed

0.2 Data Quality Analysis

A comprehensive analysis of data quality, identifying potential issues before modeling.

0.2.1 Dataset Overview

Metric	Value
Total Samples	114
Total Features	30
Numeric Features	30
Categorical Features	0
Features with Missing	0
Total Missing Values	0 (0.00%)

0.2.2 Distribution Analysis

Feature	Skewness	Kurtosis	Assessment
mean radius	1.03	1.35	Right-skewed
mean texture	0.25	-0.40	Symmetric
mean perimeter	1.07	1.56	Right-skewed
mean area	1.93	5.46	Right-skewed, Heavy-tailed

Feature	Skewness	Kurtosis	Assessment
mean smoothness	0.53	1.50	Right-skewed
mean compactness	0.91	0.54	Right-skewed
mean concavity	1.44	2.46	Right-skewed
mean concave points	1.15	1.18	Right-skewed
mean symmetry	0.52	0.34	Right-skewed
mean fractal dimension	1.26	3.76	Right-skewed, Heavy-tailed
radius error	3.30	17.46	Right-skewed, Heavy-tailed
texture error	0.75	1.02	Right-skewed
perimeter error	3.60	20.78	Right-skewed, Heavy-tailed
area error	5.87	45.73	Right-skewed, Heavy-tailed
smoothness error	1.03	1.22	Right-skewed

0.2.3 Feature Distributions

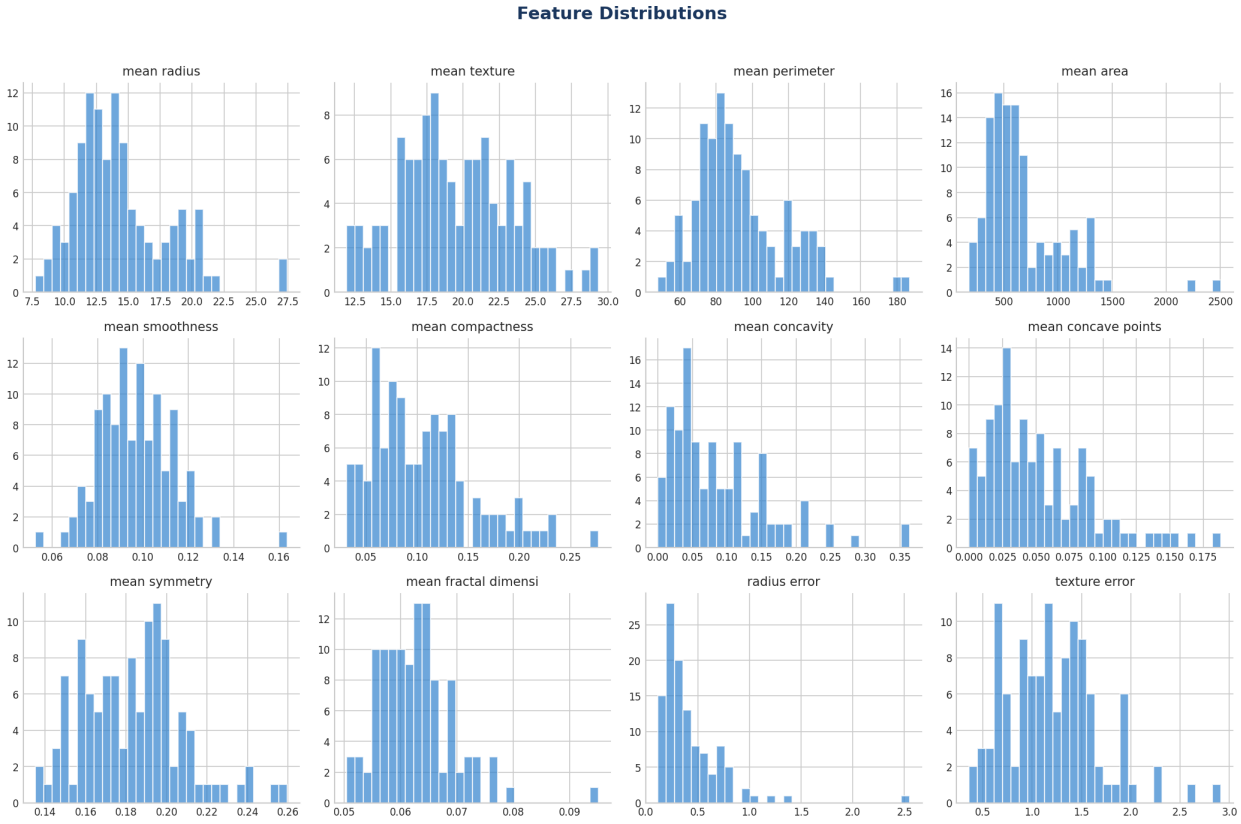

Figure 1: Feature Distributions

0.2.4 Outlier Analysis

Method: Interquartile Range (IQR) with 1.5x multiplier

Total Outliers Detected: 96 across 27 features

Feature	Outliers	% of Data	Min	Max
area error	10	8.8%	8.61	542.20
fractal dimension error	7	6.1%	0.00	0.01
concavity error	6	5.3%	0.00	0.09
mean concavity	5	4.4%	0.00	0.36
worst symmetry	5	4.4%	0.20	0.48
mean area	4	3.5%	181.00	2501.00
radius error	4	3.5%	0.12	2.55
perimeter error	4	3.5%	0.77	18.65
smoothness error	4	3.5%	0.00	0.02
symmetry error	4	3.5%	0.01	0.04
worst area	4	3.5%	268.60	4254.00
worst smoothness	4	3.5%	0.09	0.22
worst fractal dimension	4	3.5%	0.06	0.14
mean compactness	3	2.6%	0.03	0.28
mean concave points	3	2.6%	0.00	0.19

0.2.5 Outlier Distribution

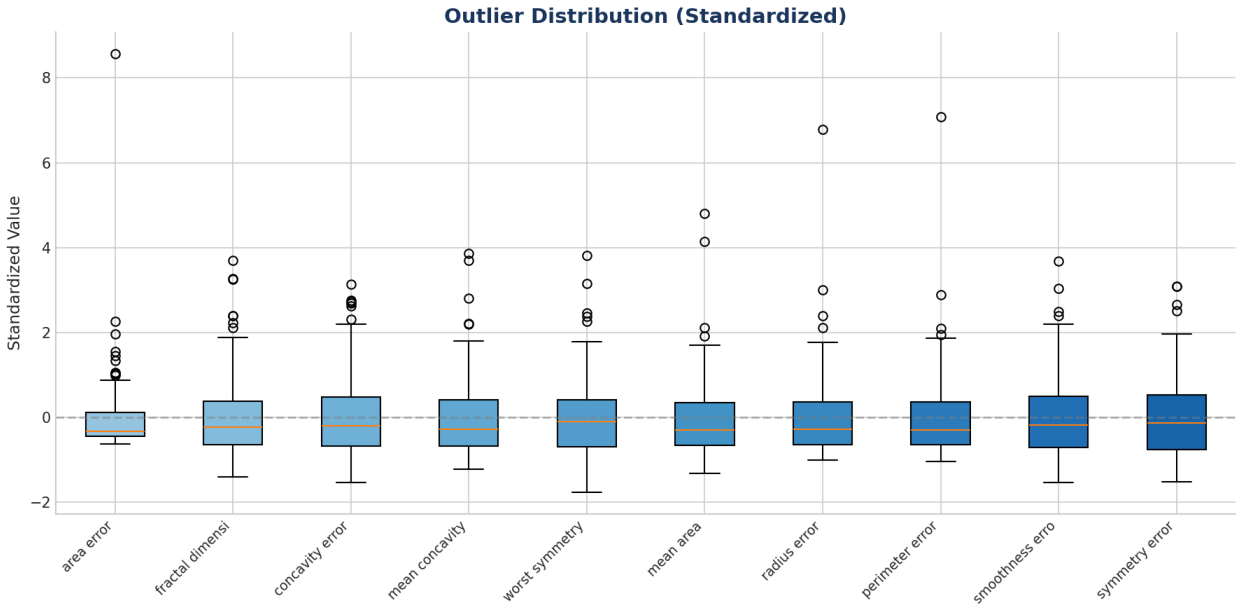

Figure 2: Outlier Boxplot

0.3 Correlation & Multicollinearity Analysis

Understanding feature relationships is critical for model interpretation and feature selection.

0.3.1 Correlation Matrix

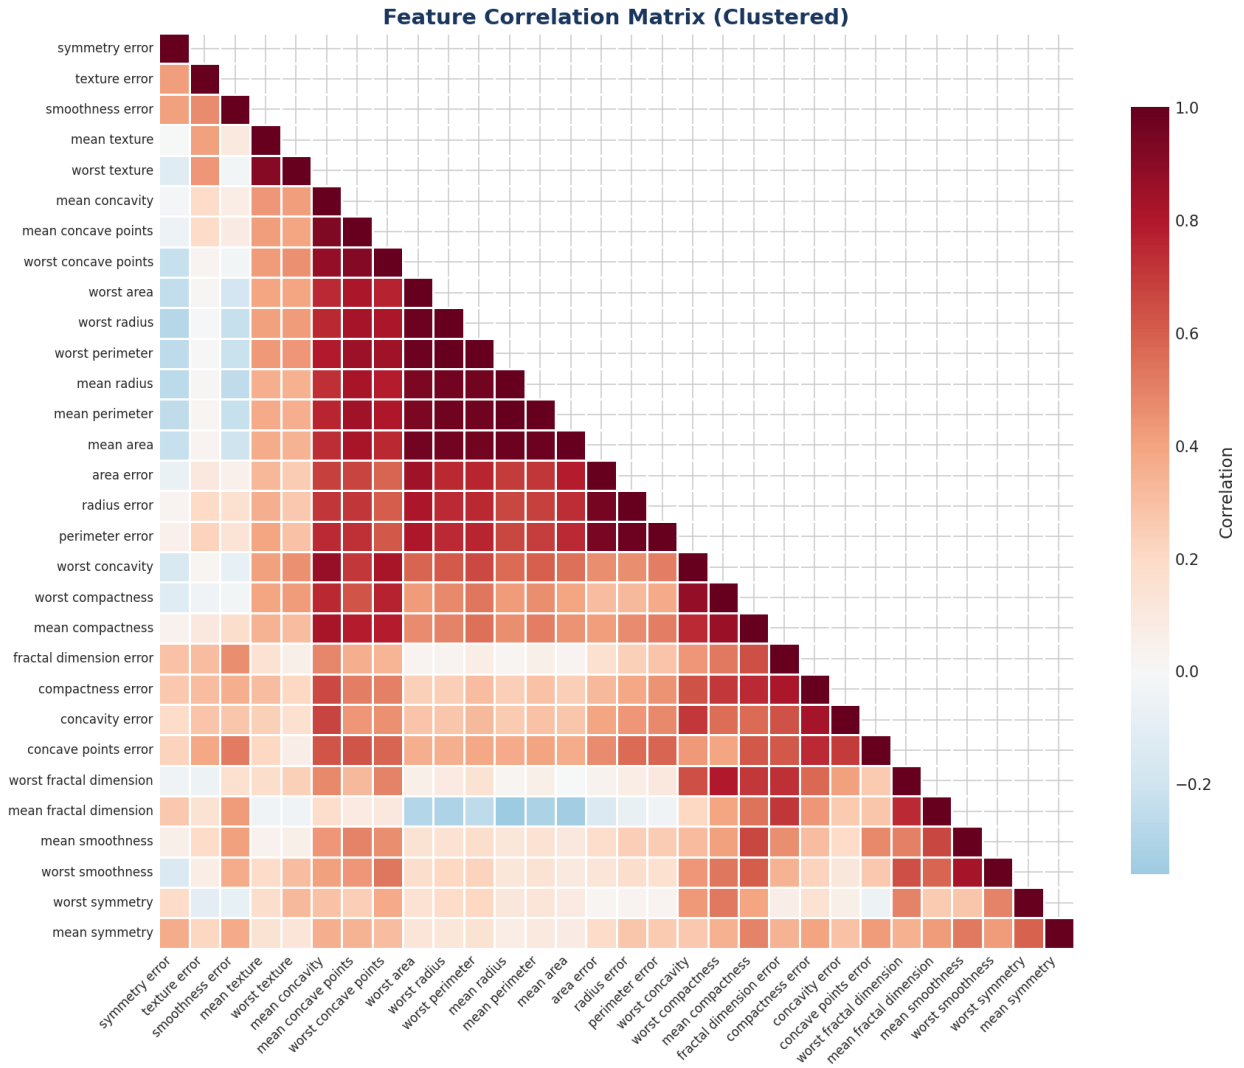

Figure 3: Correlation Matrix

0.3.2 Highly Correlated Feature Pairs (|r| >= 0.7)

Feature 1	Feature 2	Correlation
mean radius	mean perimeter	0.998
worst radius	worst perimeter	0.994
mean perimeter	mean area	0.983

Feature 1	Feature 2	Correlation
mean radius	mean area	0.983
worst radius	worst area	0.980
radius error	perimeter error	0.979
worst perimeter	worst area	0.976
mean perimeter	worst radius	0.969
mean perimeter	worst perimeter	0.969
mean radius	worst radius	0.968
mean area	worst area	0.966
mean area	worst radius	0.963
mean radius	worst perimeter	0.962
mean area	worst perimeter	0.959
radius error	area error	0.953

0.3.3 Variance Inflation Factor (VIF)

VIF measures multicollinearity. VIF > 5 indicates moderate, VIF > 10 indicates severe multicollinearity.

Feature	VIF	Assessment
mean radius	79262.63	Critical
mean perimeter	72990.12	Critical
worst radius	22342.62	Critical
worst perimeter	12079.97	Critical
worst area	2216.27	Critical
mean area	1946.14	Critical
worst fractal dimension	1185.56	Critical
mean fractal dimension	1108.49	Critical
worst smoothness	849.75	Critical
mean smoothness	811.67	Critical
worst texture	667.77	Critical
radius error	547.22	Critical
perimeter error	478.85	Critical
mean texture	478.71	Critical
worst symmetry	319.25	Critical

0.3.4 VIF Visualization

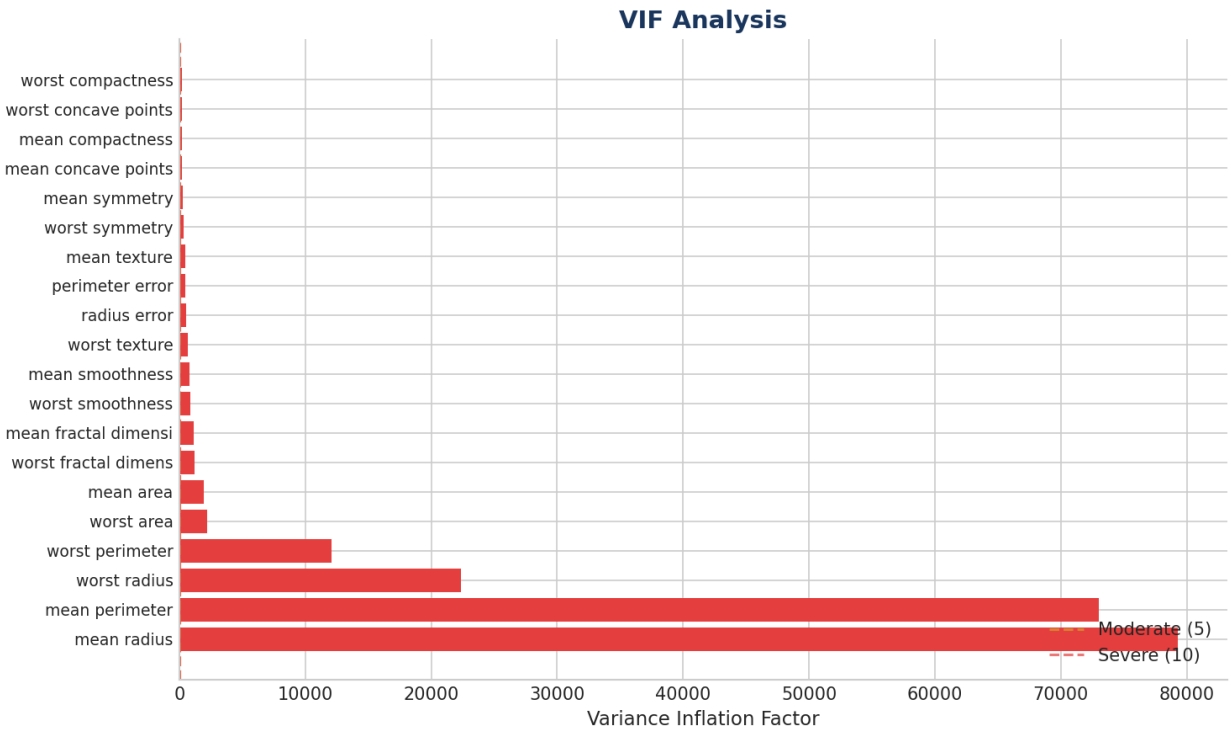

Figure 4: VIF Analysis

0.4 Data Profile

0.4.1 Dataset Overview

- **Total Samples:** 569
- **Total Features:** 30

0.4.2 Data Types

Data Type	Count
float64	30

0.4.3 EDA Recommendations

- All features are numeric - no encoding needed
- Features describe cell nuclei measurements
- Consider feature scaling for distance-based models

0.5 Feature Importance Analysis

Feature importance measures how much each feature contributes to the model's predictions. Higher values indicate more influential features.

0.5.1 Top 20 Features

Rank	Feature	Importance
1	worst texture	1.2551
2	radius error	1.0830
3	worst concave points	0.9537
4	worst area	0.9478
5	worst radius	0.9476
6	worst symmetry	0.9392
7	area error	0.9291
8	worst concavity	0.8232
9	worst perimeter	0.7632
10	worst smoothness	0.7466
11	mean concave points	0.7042
12	mean compactness	0.6483
13	compactness error	0.6472
14	mean concavity	0.6021
15	mean texture	0.5527
16	perimeter error	0.5443
17	mean area	0.5411
18	mean radius	0.5115
19	mean perimeter	0.4763
20	concave points error	0.4438

0.5.2 Feature Importance Visualization

0.6 Model Benchmarking

Multiple machine learning algorithms were evaluated using 5-fold cross-validation. The table below shows the performance of each model.

0.6.1 Model Comparison

Model	CV Score	Std Dev	Test Score	Time (s)
Logistic Regression	0.9802	±0.0128	0.9825	0.53
Random Forest	0.9538	±0.0235	0.9561	0.17
Gradient Boosting	0.9560	±0.0139	0.9561	0.49

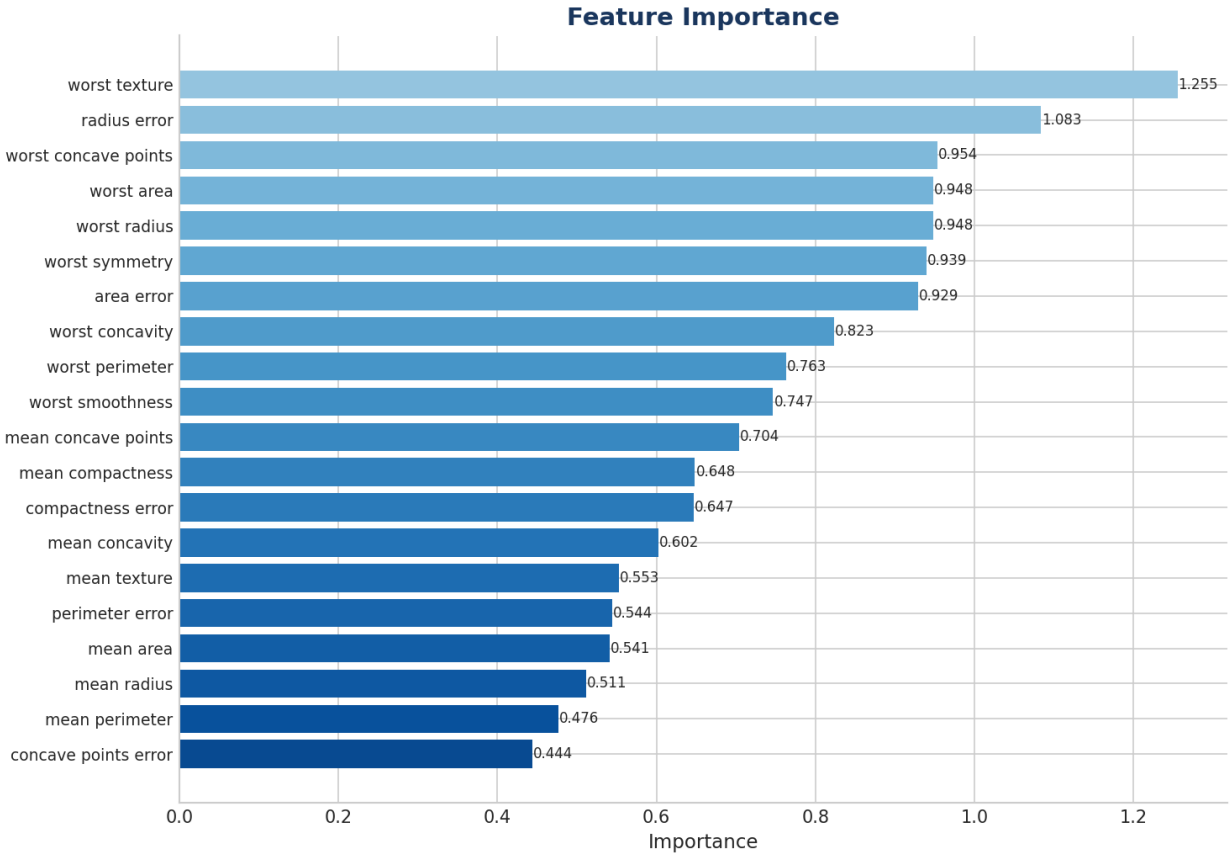

Figure 5: Feature Importance

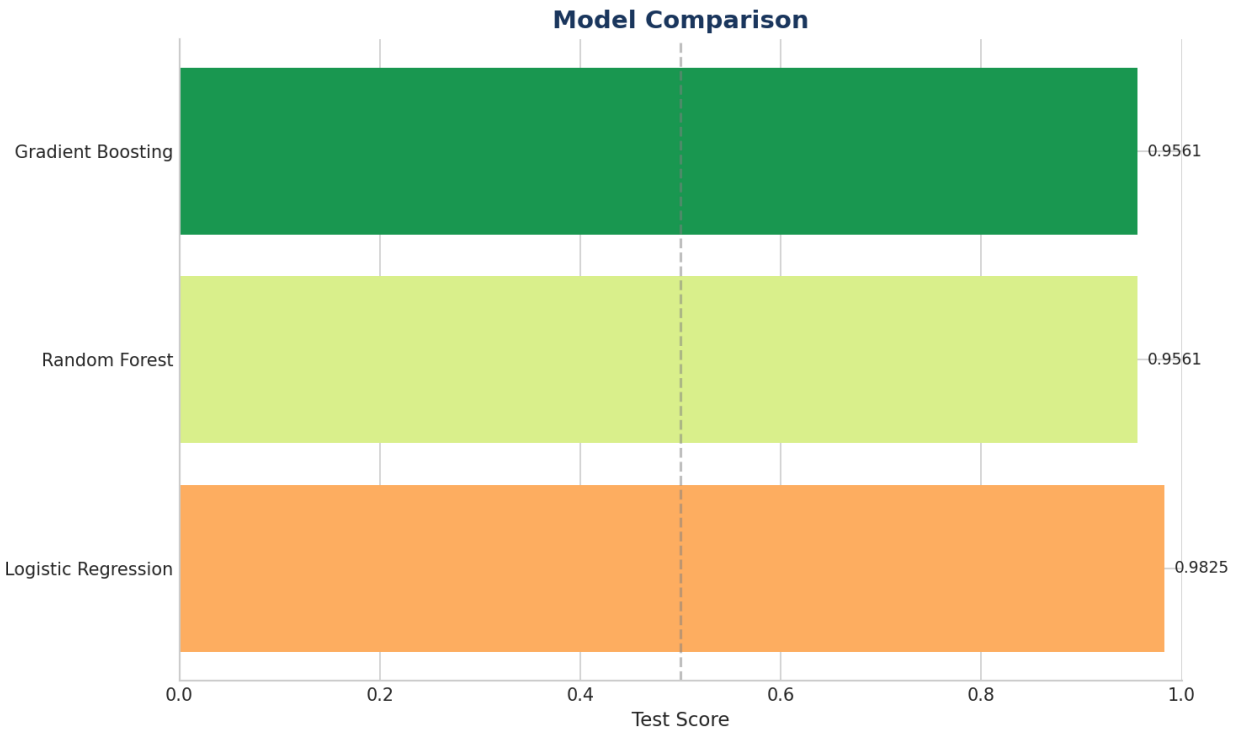

Figure 6: Model Benchmarking

0.6.2 Performance Visualization

0.7 Learning Curve Analysis

Learning curves reveal how model performance changes with training data size, helping diagnose underfitting vs overfitting.

0.7.1 Bias-Variance Diagnosis

Good Fit: Model has balanced bias-variance tradeoff.

Metric	Value
Final Training Score	0.9736
Final Validation Score	0.9391
Gap (Train - Val)	0.0345
Training Samples Used	91

0.7.2 Learning Curve Visualization

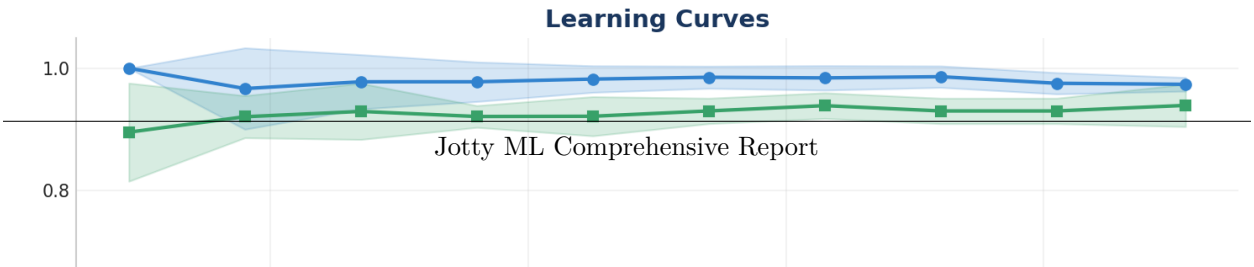

0.7.3 Interpretation Guide

- **Converging curves** with small gap → Good fit
- **Flat training curve** at low score → High bias, need more complex model
- **Large gap** between curves → High variance, need regularization or more data
- **Curves still improving** → May benefit from more training data

---

0.8 Cross-Validation Detailed Analysis

5-fold cross-validation provides robust performance estimates and helps detect instability.

0.8.1 Fold-by-Fold Results

Fold	Train Accuracy	Test Accuracy	Train F1	Test F1
1	0.9780	0.9130	0.9762	0.9115
2	0.9670	0.9565	0.9641	0.9533
3	0.9890	0.9130	0.9883	0.9042
4	0.9780	0.9130	0.9762	0.9042
5	0.9565	1.0000	0.9527	1.0000

0.8.2 Stability Analysis

Metric	Value
Mean Accuracy	0.9391
Std Deviation	0.0348
CV Coefficient	3.70%
95% CI	[0.8710, 1.0073]

Stability Assessment: Good

0.8.3 CV Performance Distribution

0.9 Classification Performance

0.9.1 Classification Report

Class	Precision	Recall	F1-Score	Support
Malignant	0.976	0.976	0.976	42
Benign	0.986	0.986	0.986	72

Class	Precision	Recall	F1-Score	Support
Accuracy			**0.982**	

0.9.2 Confusion Matrix

0.10 ROC Curve Analysis

The Receiver Operating Characteristic (ROC) curve shows the trade-off between true positive rate and false positive rate at various classification thresholds.

0.10.1 Key Metrics

- **AUC-ROC:** 0.9954
- **Optimal Threshold:** 0.3659

0.10.2 ROC Curve

0.11 Precision-Recall Analysis

The Precision-Recall curve is especially useful for imbalanced datasets, showing the trade-off between precision and recall.

0.11.1 Key Metrics

- **Average Precision:** 0.9971

0.11.2 Precision-Recall Curve

0.12 Probability Calibration Analysis

Well-calibrated probabilities are essential for reliable decision-making. A perfectly calibrated model's predicted probabilities should match actual outcome frequencies.

0.12.1 Calibration Metrics

Metric	Value	Interpretation
Brier Score	0.0215	Lower is better (0 = perfect)
Expected Calibration Error	0.0222	Lower is better

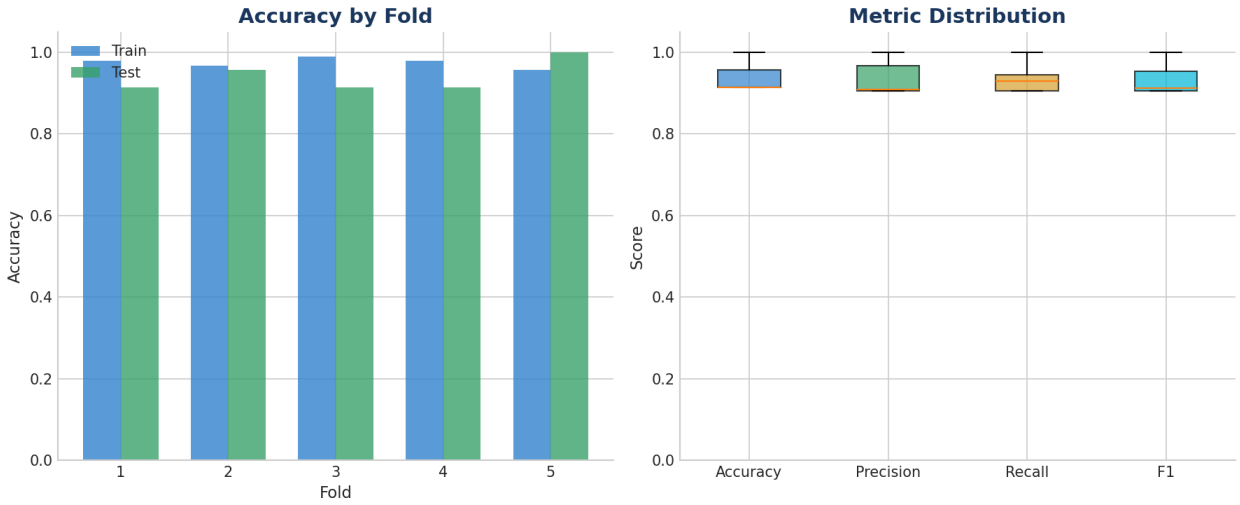

Figure 8: CV Analysis

0.12.2 Calibration Curve

0.12.3 Interpretation

- Points on diagonal = perfectly calibrated
- Points above diagonal = underconfident (probabilities too low)
- Points below diagonal = overconfident (probabilities too high)

0.13 Lift & Gain Analysis

These charts help evaluate model effectiveness for targeted campaigns and prioritization.

0.13.1 Key Metrics

Metric	Value	Interpretation
KS Statistic	0.3596	Maximum separation between model and random
KS at Decile	64%	Optimal cutoff point
Top 10% Lift	1.58x	Model advantage in top 10%
Top 20% Lift	1.58x	Model advantage in top 20%

0.13.2 Cumulative Gains & Lift Curves

0.13.3 Business Interpretation

- **Gains Curve**: Shows % of positives captured by targeting top X% of predictions
- **Lift Curve**: Shows how much better the model is vs random selection
- **KS Statistic**: Higher values indicate better model discrimination

---

0.14 Threshold Optimization

Choosing the right classification threshold depends on business objectives.

0.14.1 Optimal Thresholds

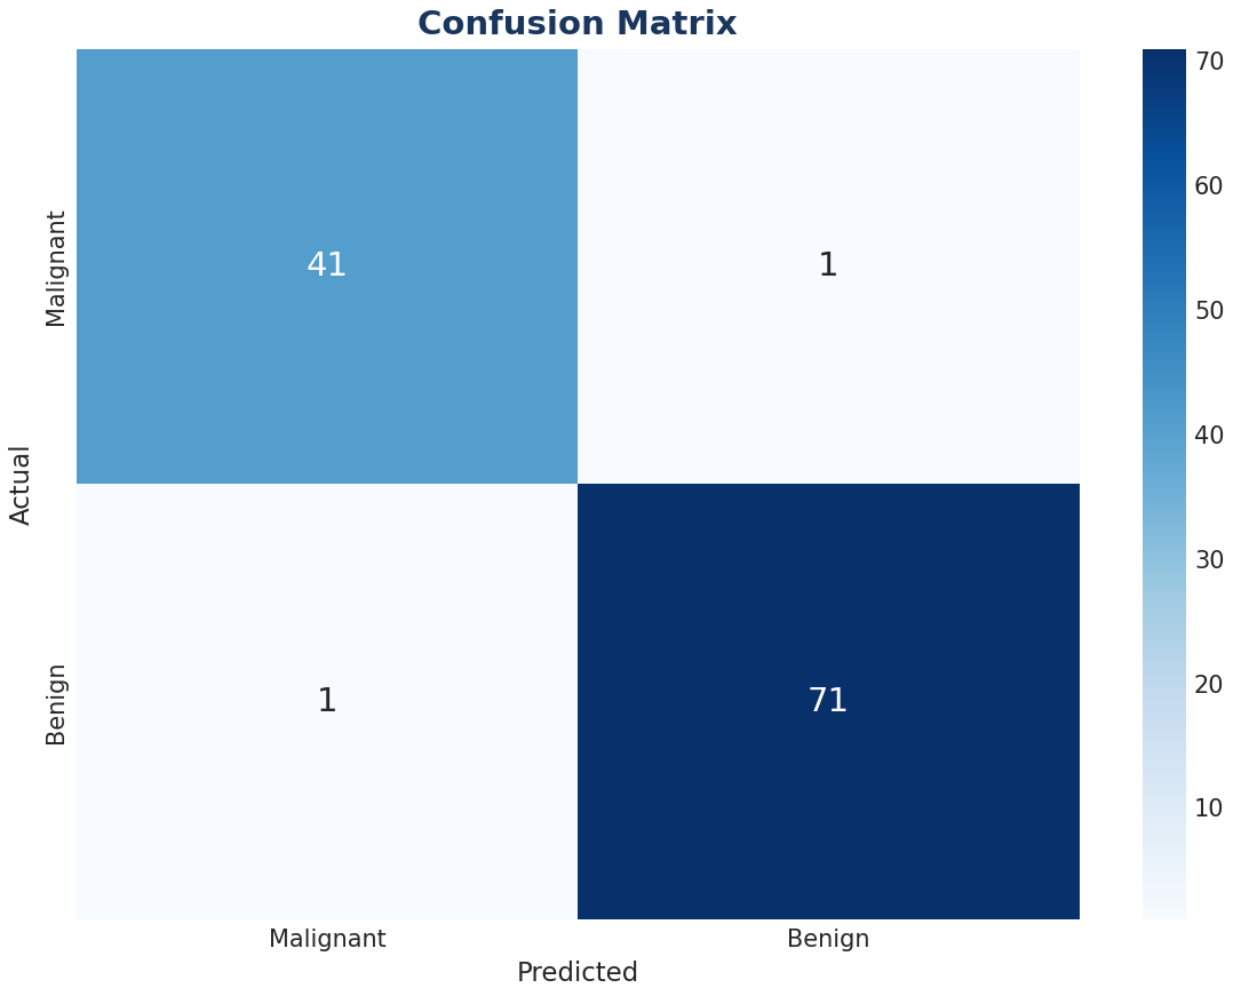

Figure 9: Confusion Matrix

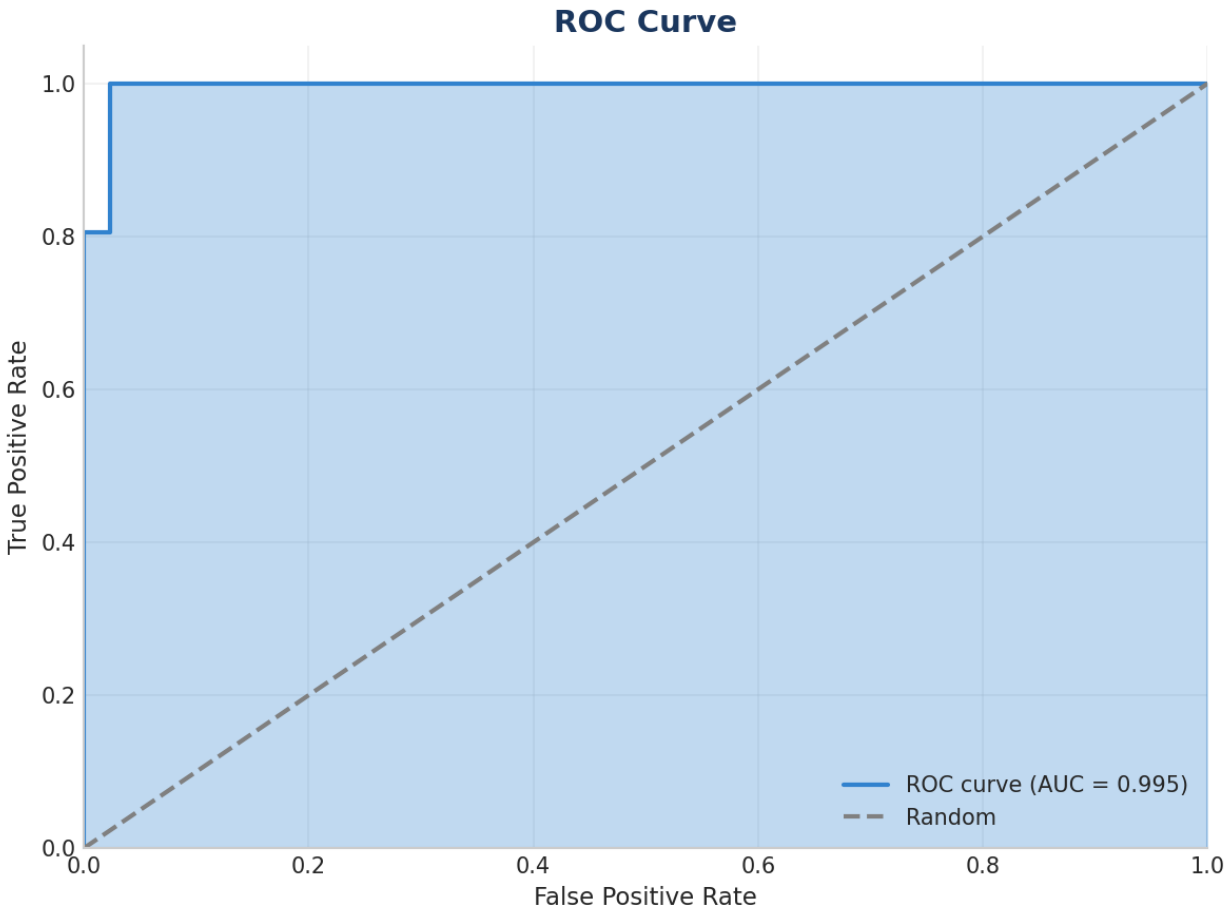

Figure 10: ROC Curve

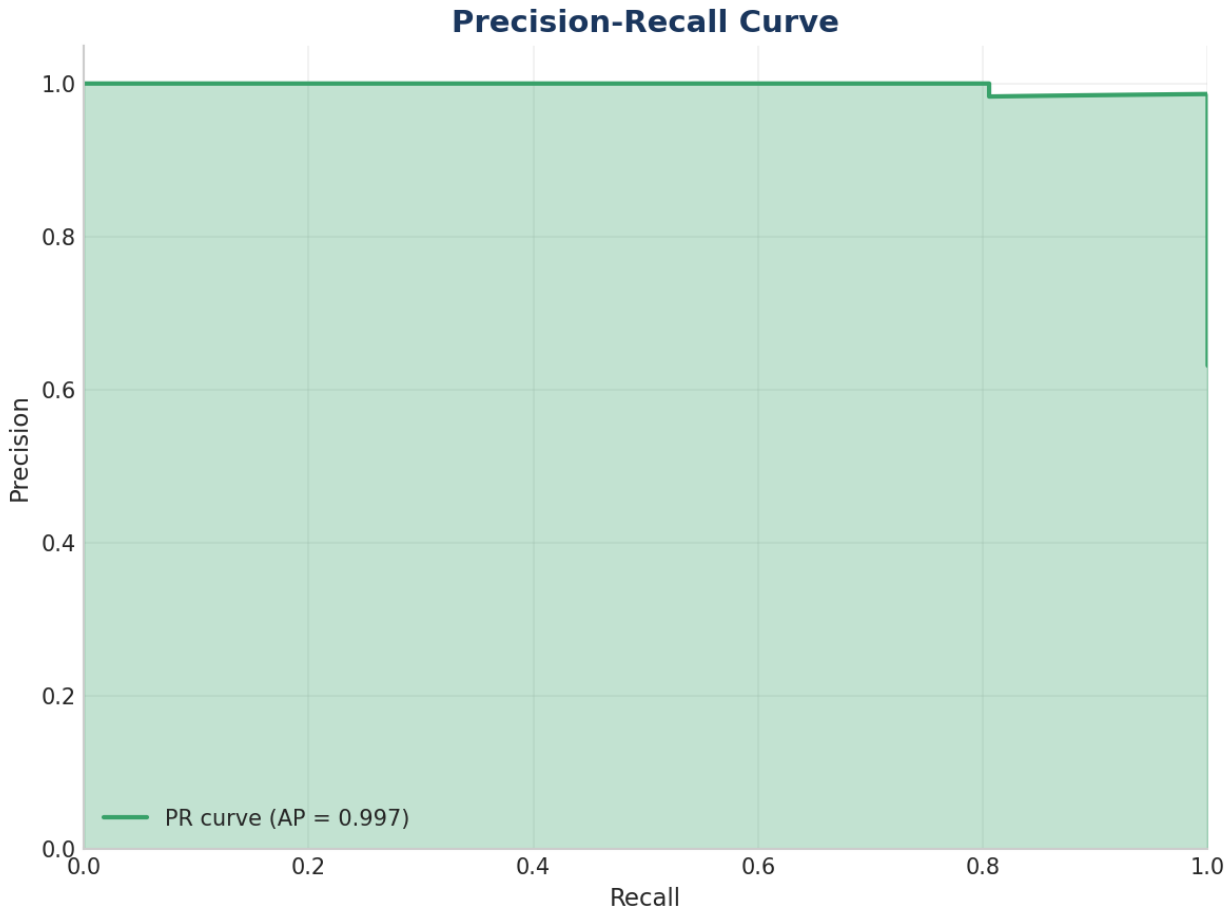

Figure 11: Precision-Recall Curve

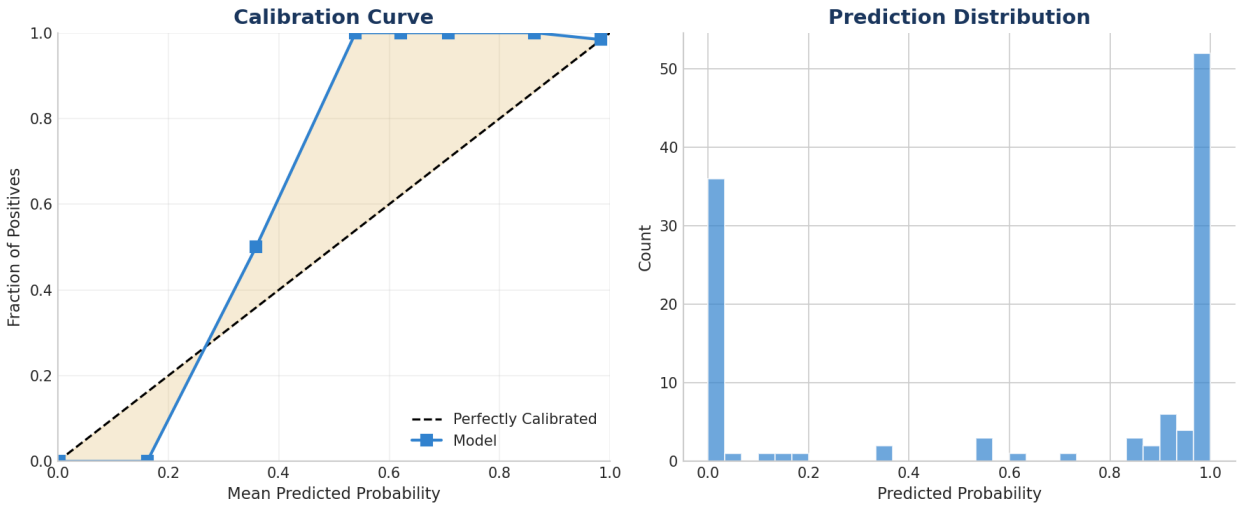

Figure 12: Calibration Curve

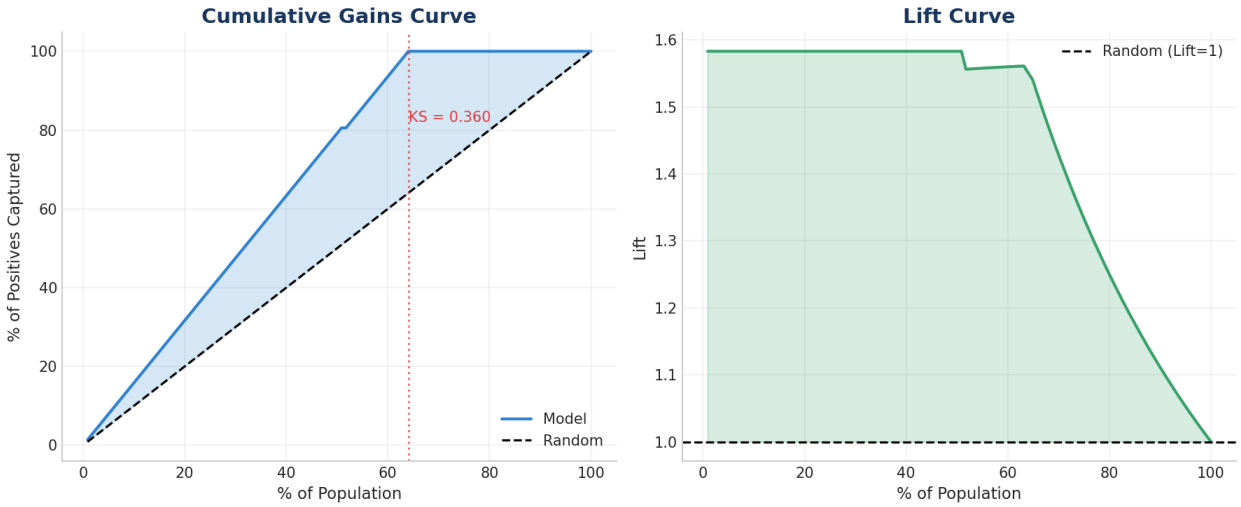

Figure 13: Lift and Gain Charts

Threshold	TP	FP	FN	TN	Precision	Recall	F1
0.10	72	5	0	37	0.935	1.000	0.966
0.20	72	2	0	40	0.973	1.000	0.986
0.30	72	2	0	40	0.973	1.000	0.986
0.40	71	1	1	41	0.986	0.986	0.986
0.50	71	1	1	41	0.986	0.986	0.986
0.60	68	1	4	41	0.986	0.944	0.965
0.70	67	1	5	41	0.985	0.931	0.957
0.80	66	1	6	41	0.985	0.917	0.950
0.90	61	1	11	41	0.984	0.847	0.910

0.14.3 Threshold Visualization

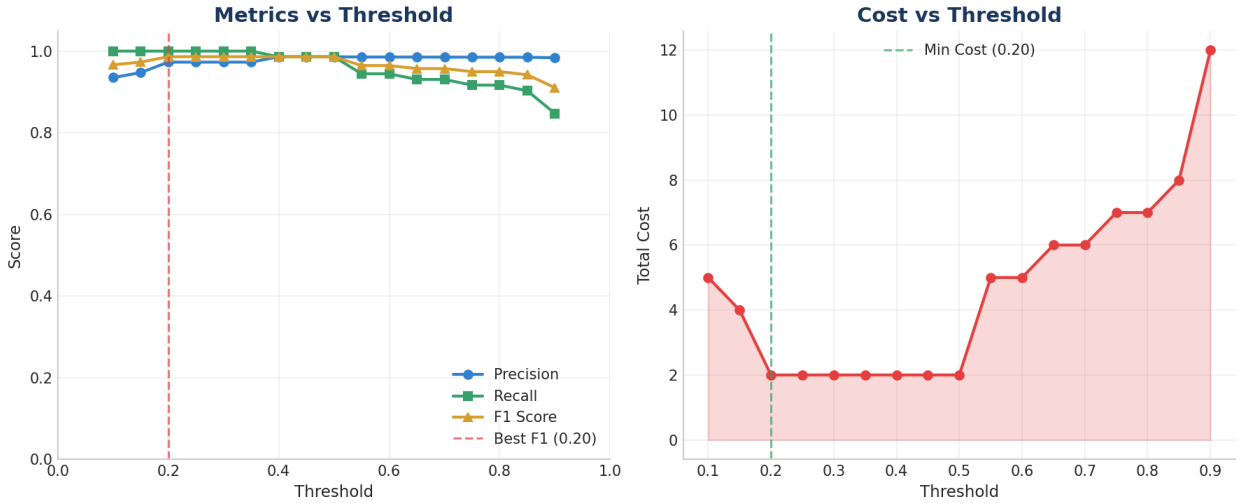

Figure 14: Threshold Analysis

0.14.4 Cost Parameters Used

Metric	Value
Total Errors	2
Error Rate	1.75%
Accuracy	98.25%

0.15.2 Confusion Matrix Breakdown

- Class 0 misclassified as Class 1: 1 (2.4%)
- Class 1 misclassified as Class 0: 1 (1.4%)

0.15.3 Hardest to Classify Samples (Most Confident Errors)

Sample	True	Predicted	Probability	Confidence
53	0	1	0.909	0.817
16	1	0	0.366	0.268

0.15.4 Error Distribution Analysis

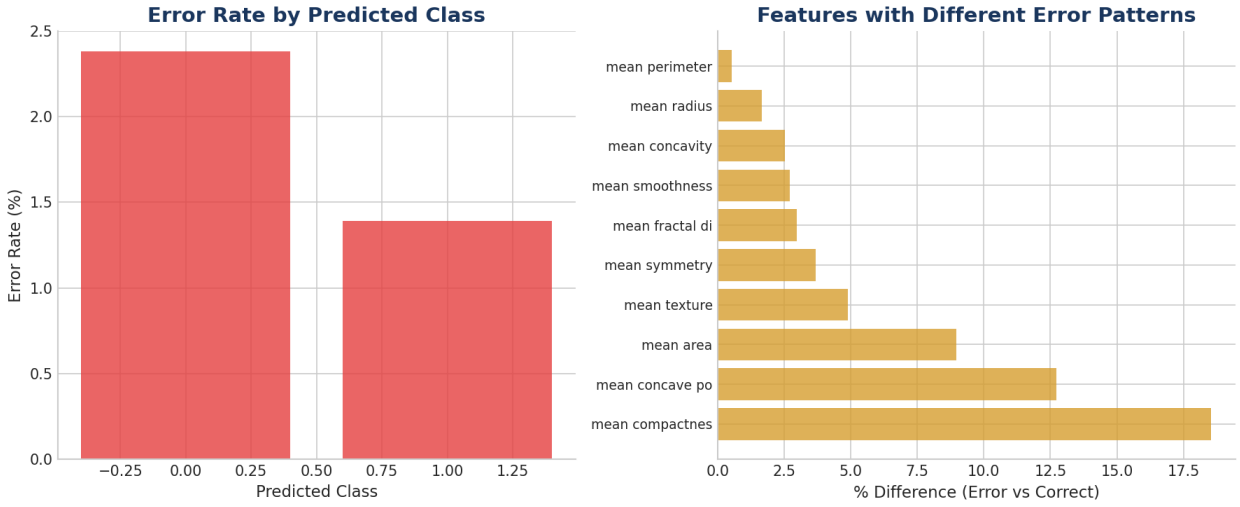

Figure 15: Error Analysis

0.16 SHAP Deep Analysis

SHAP (SHapley Additive exPlanations) provides consistent, locally accurate feature attributions for any machine learning model.

0.16.1 Global Feature Importance (Mean |SHAP|)

Rank	Feature	Mean	SHAP
1	worst texture	0.9617	7.8%
2	radius error	0.7973	14.2%
3	worst radius	0.7614	20.3%
4	worst area	0.7453	26.3%
5	worst concave points	0.7220	32.1%
6	area error	0.6391	37.3%
7	worst perimeter	0.6109	42.2%
8	worst smoothness	0.6033	47.1%
9	worst concavity	0.5463	51.5%
10	worst symmetry	0.5358	55.8%
11	mean concave points	0.5358	60.1%
12	mean compactness	0.4774	64.0%
13	compactness error	0.4310	67.4%
14	mean area	0.4174	70.8%
15	mean concavity	0.3991	74.0%

0.16.2 SHAP Summary Plot

Shows feature impact on predictions. Color indicates feature value (red=high, blue=low).

0.16.3 SHAP Feature Importance Bar

0.16.4 SHAP Dependence Plots (Top 3 Features)

Shows how feature values affect SHAP values, revealing non-linear relationships.

0.16.5 SHAP Waterfall (Sample Prediction)

Shows how features contribute to a single prediction.

0.17 Baseline Comparison

0.17.1 Performance Improvement

Model	Score	Improvement
Baseline	0.6270	-
Best Model	**0.9825**	**+0.3555 (+56.7%)**

The final model achieves a **56.7%** improvement over the baseline.

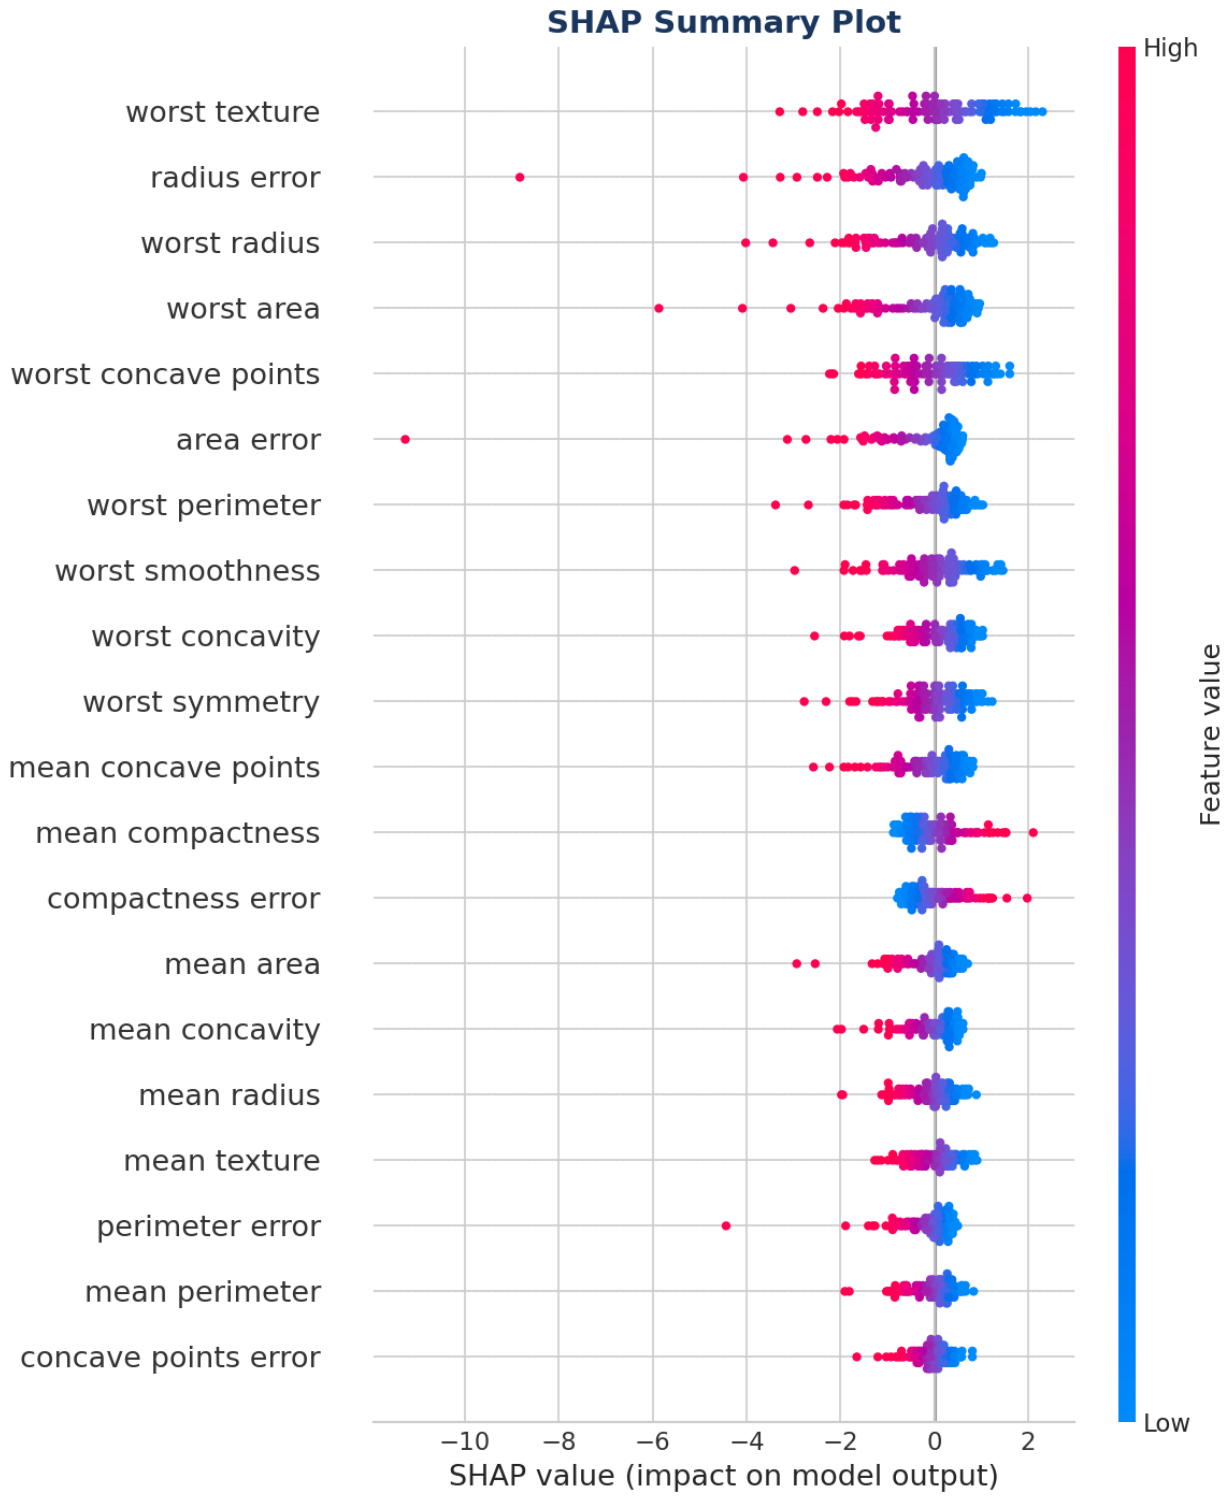

Figure 16: SHAP Summary

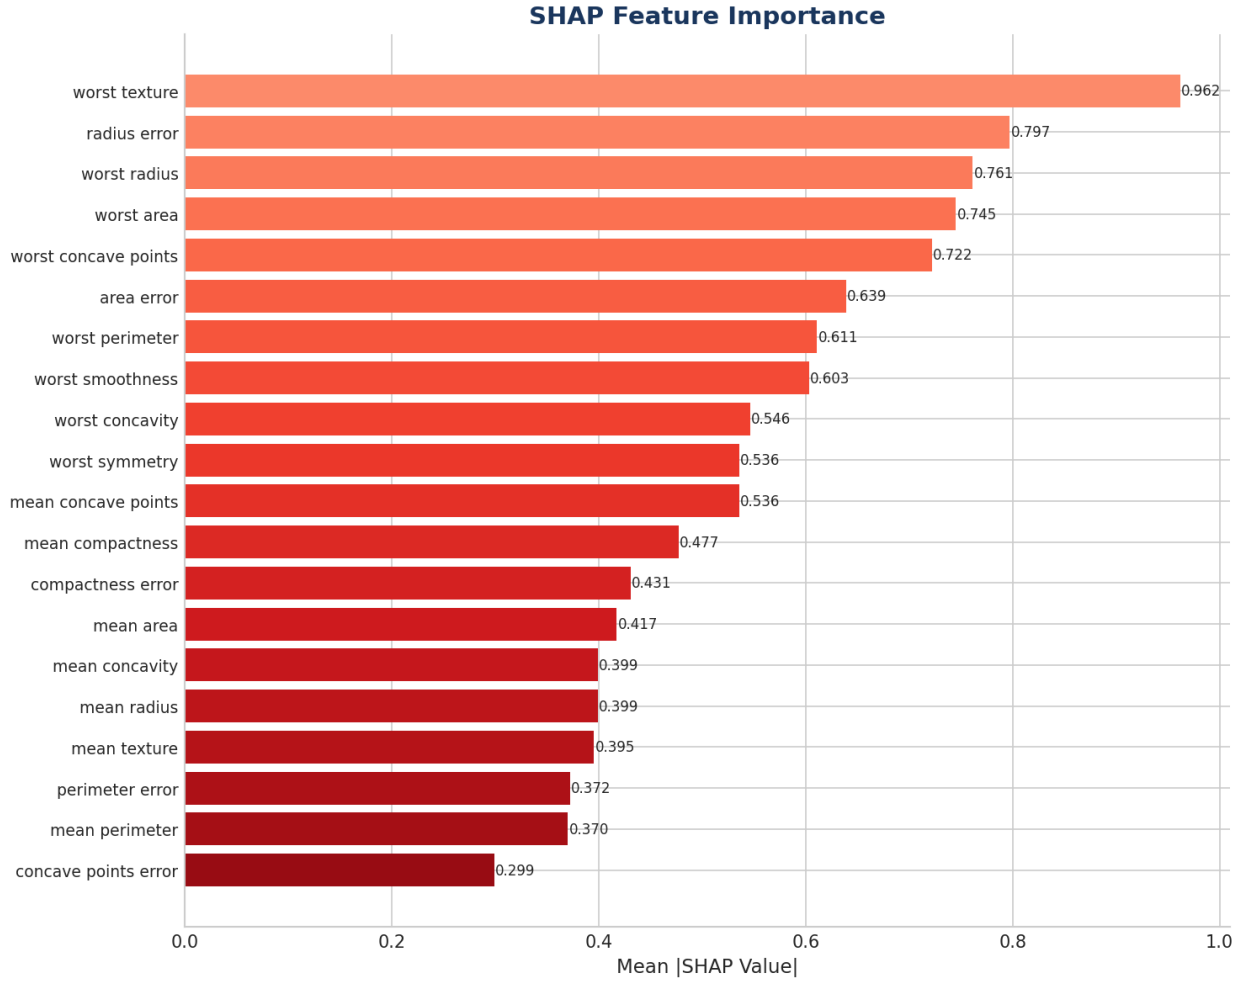

Figure 17: SHAP Bar Plot

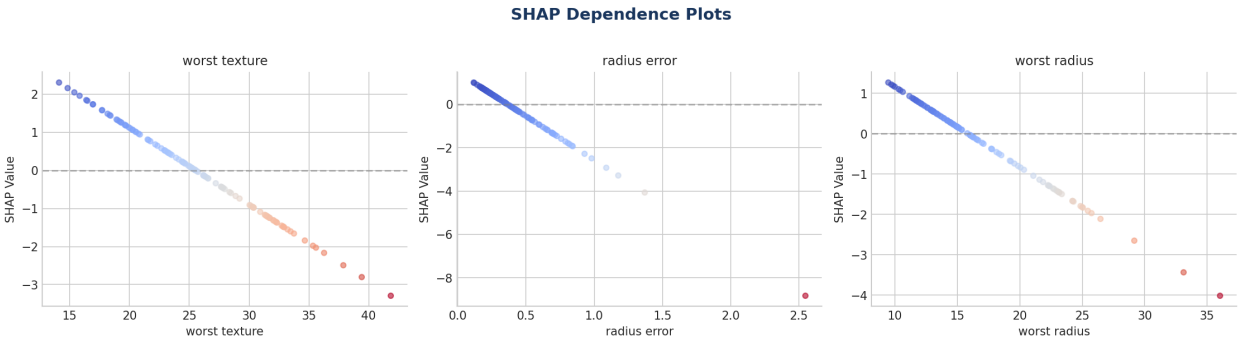

Figure 18: SHAP Dependence

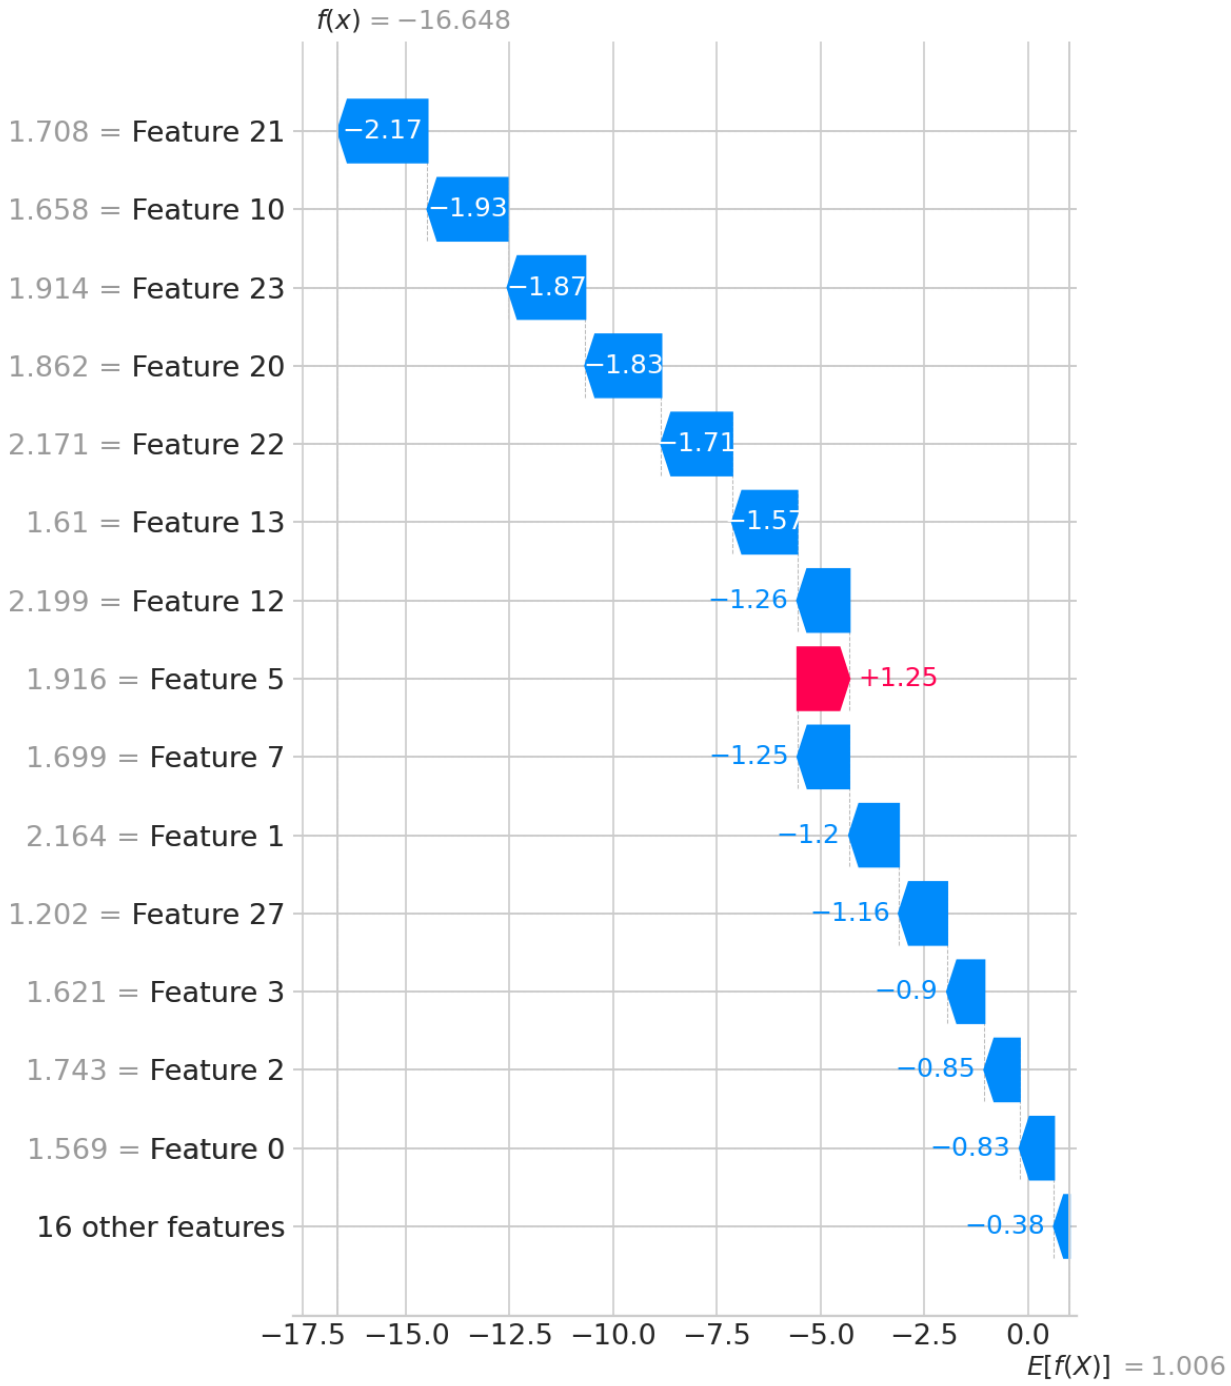

Figure 19: SHAP Waterfall

0.18 Recommendations & Next Steps

1. Excellent performance (98.2%) achieved - monitor for overfitting
2. Logistic Regression provides good interpretability - ideal for regulated industries
3. Top predictive features: worst texture, radius error, worst concave points
4. Excellent discrimination (AUC=0.995) - suitable for production
5. Monitor model performance over time for concept drift
6. Validate on held-out data before production deployment
7. Document model decisions for regulatory compliance

Report generated by Jotty SwarmMLComprehensive on 2026-02-05 03:18:31

0.19 Reproducibility

Full information for reproducing this analysis.

0.19.1 Model Configuration

Model Type: LogisticRegression

Hyperparameters:

Parameter	Value

0.19.2 Random Seeds

Component	Seed
Main Random State	42
NumPy	42
Train/Test Split	42

0.19.3 Environment

Component	Version
Python Version	3.11.2
Platform	Linux-5.4.17-2136.312.3.4.el8uek.aarch64-aarch64-with-glibc2.36
Processor	

0.19.4 Package Versions

Package	Version
numpy	1.26.4
pandas	2.3.3
matplotlib	3.10.8
seaborn	0.13.2
shap	0.50.0

0.19.5 Generation Timestamp

Report Generated: 2026-02-05 03:18:31
